



HOW TO BUILD A CHATBOT

Session 5 -
Building a Chatbot

SESSION 5

AGENDA



1

Demo of Target Solution

2

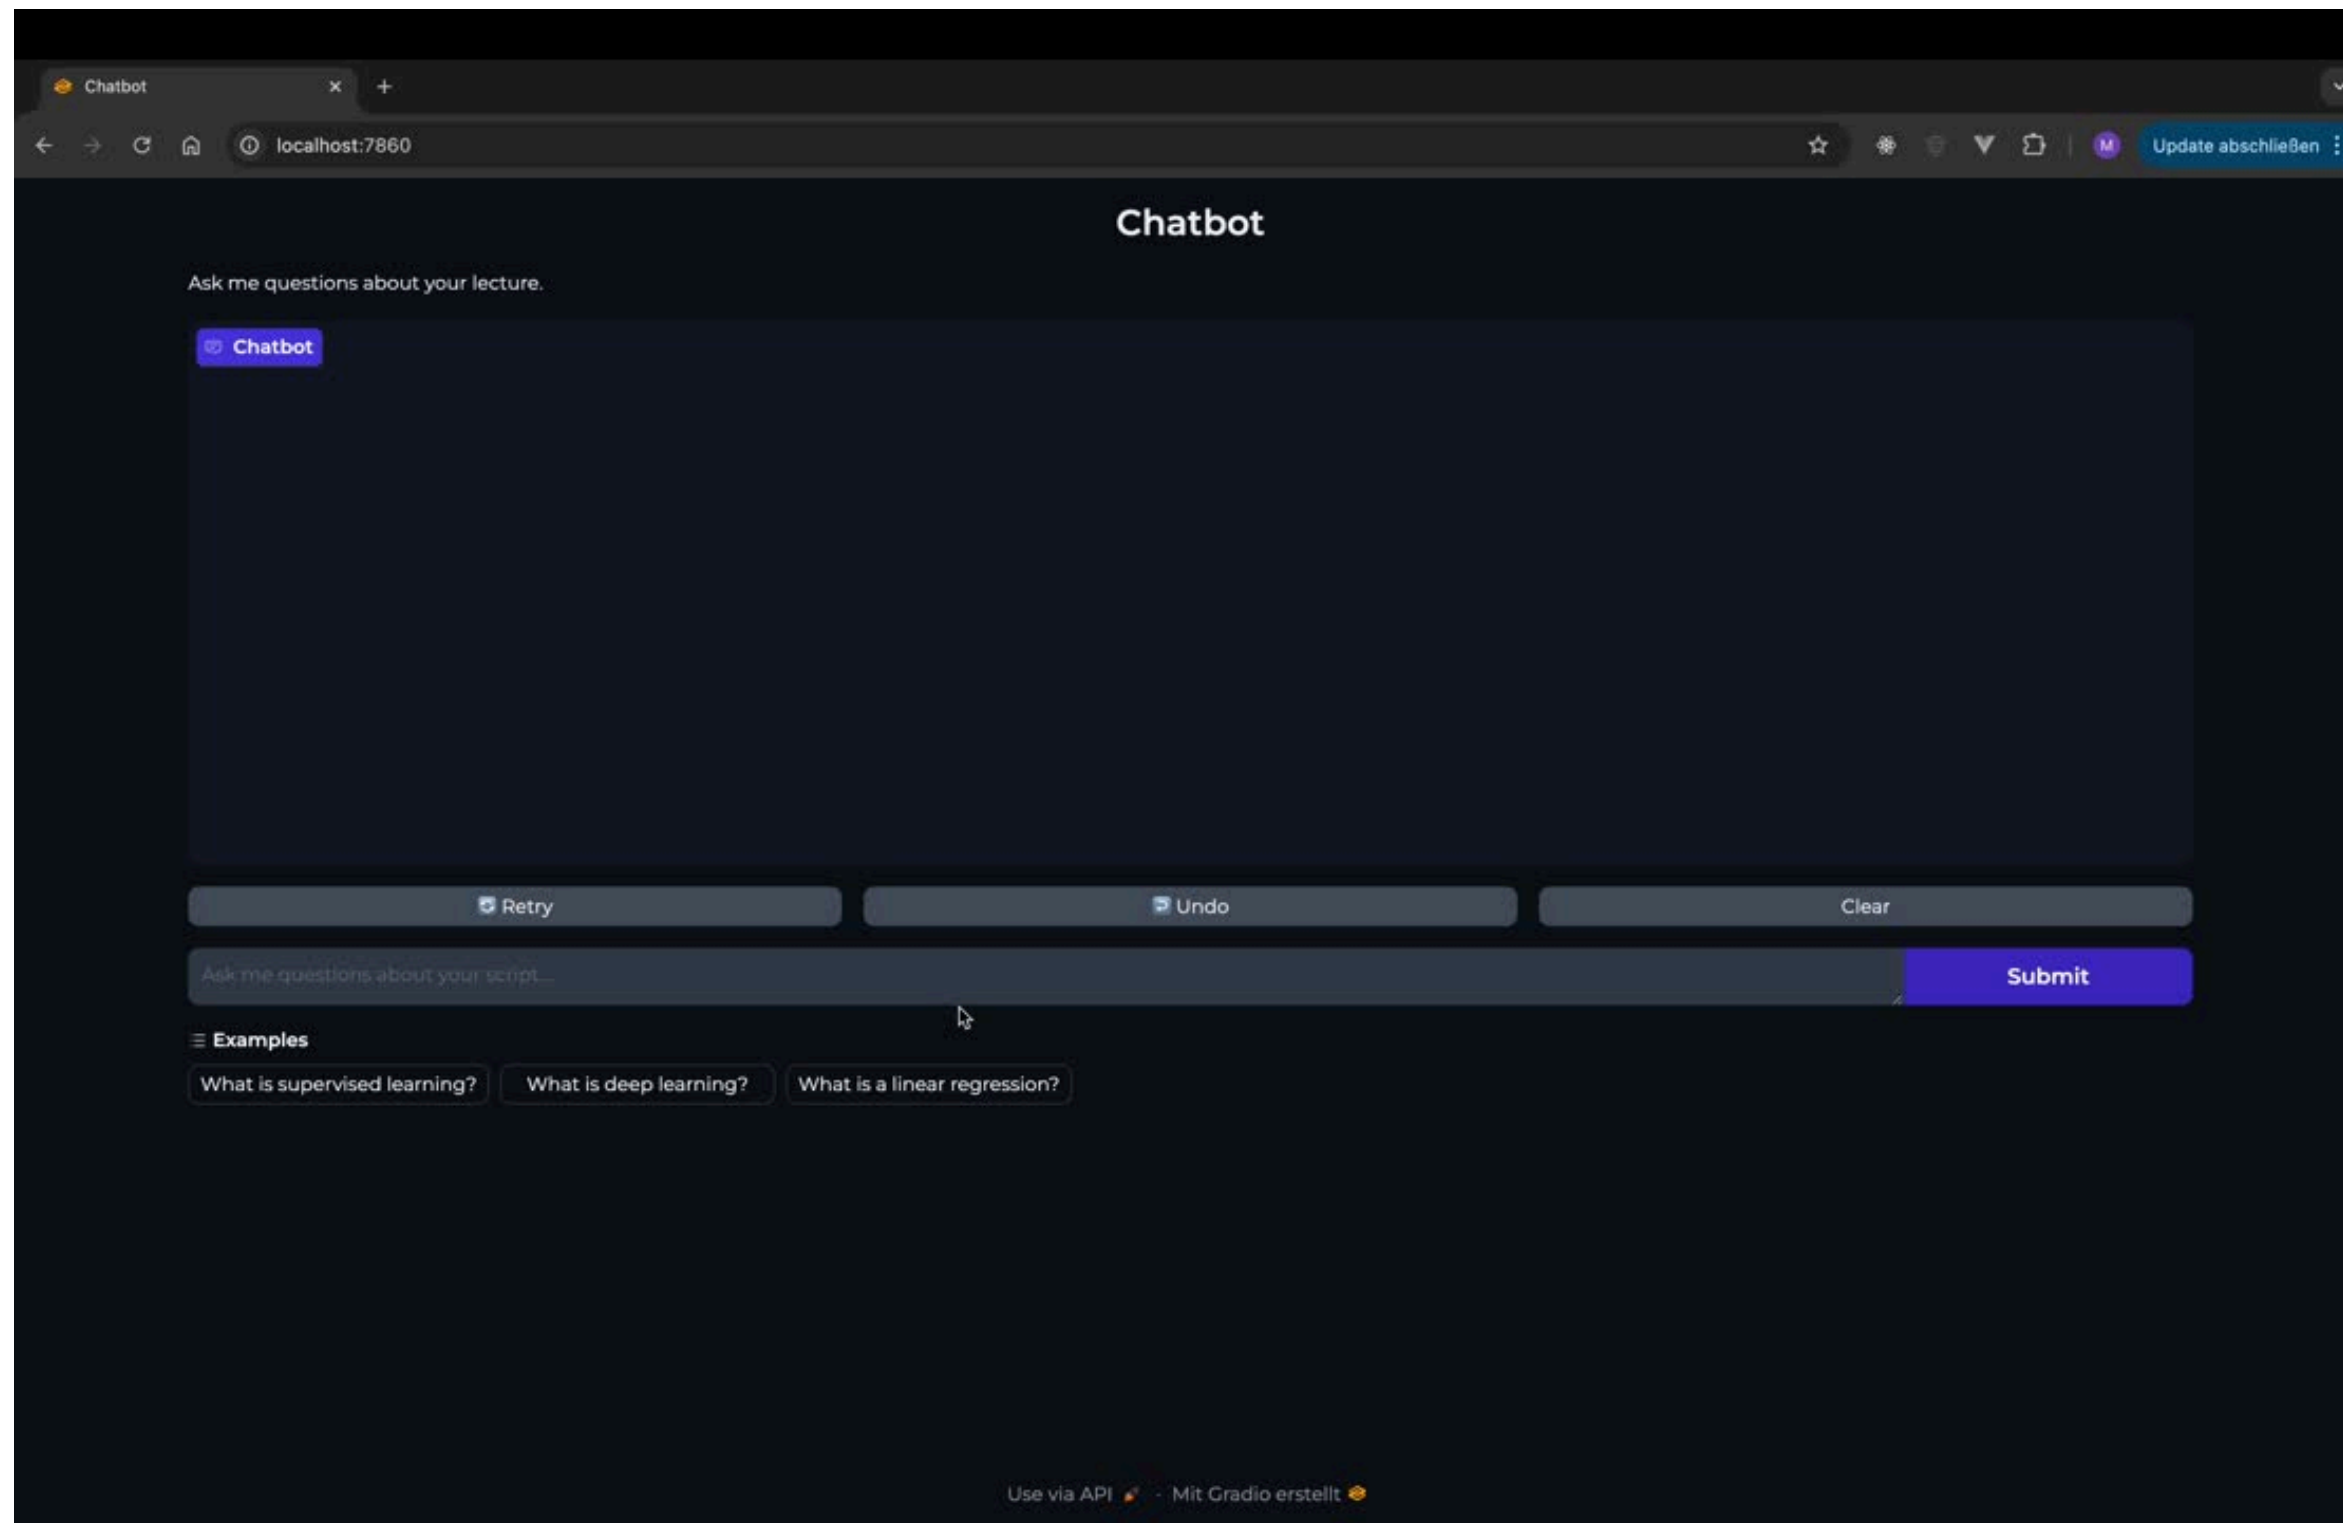
Target Architecture

3

Building Blocks

DEMO OF TARGET SOLUTION

Chatbot app in action.



TARGET ARCHITECTURE

Frontend:

- Web app built with Gradio, accessible via browser.

Backend:

- Python-based with FastAPI and LangChain.

LLM Serving:


- Ollama for managing large language models.

Knowledge Storage:

- Vector database for knowledge management.

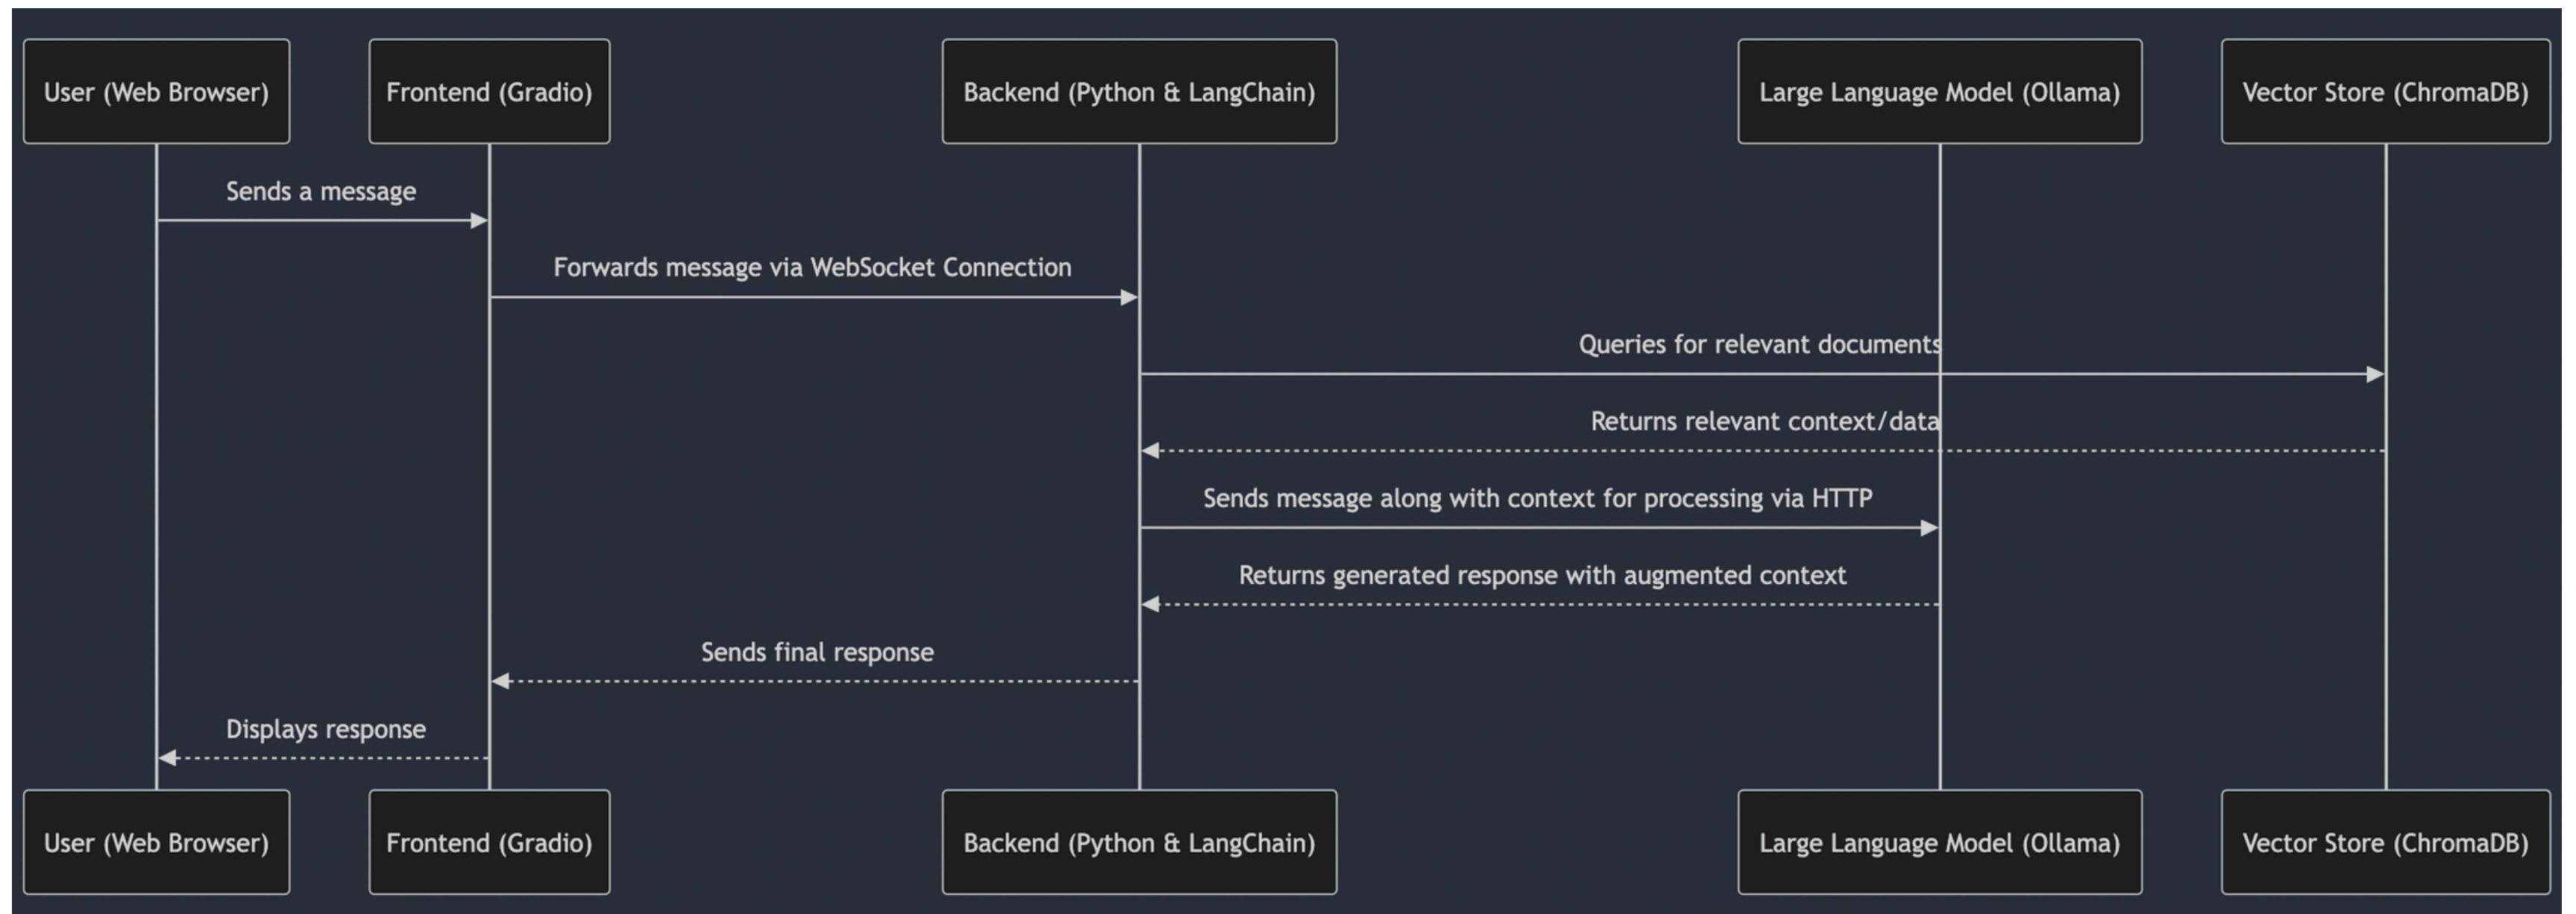
Deployment:

- Docker containers for application deployment.



TARGET ARCHITECTURE

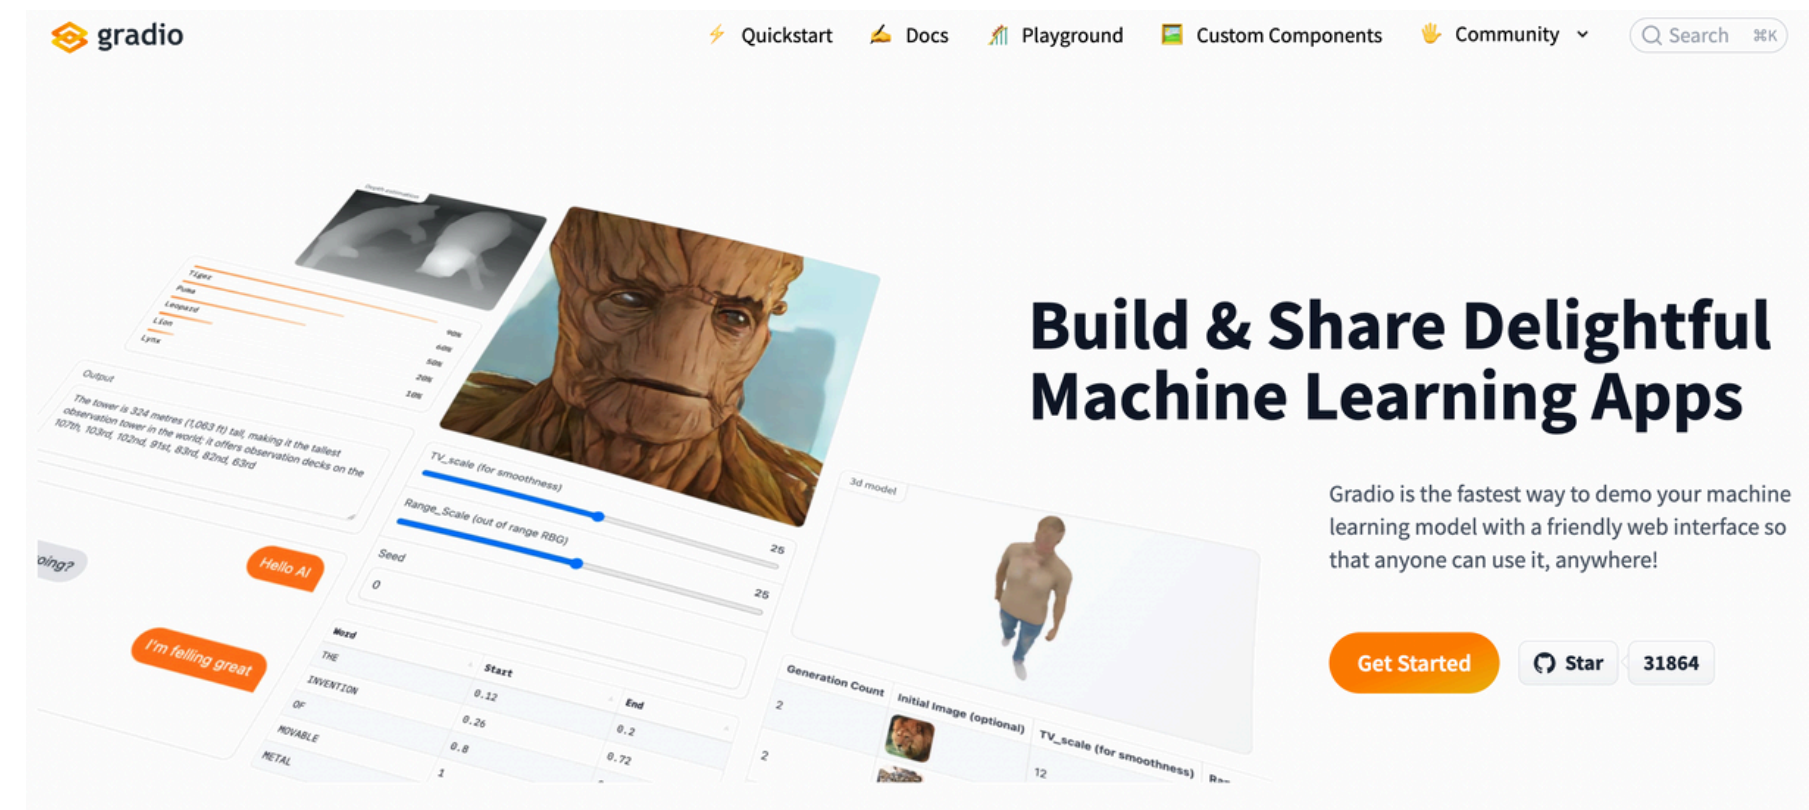
User interaction workflow.



BUILDING BLOCKS

Gradio Frontend (webapp)

- Open-source Python library
- Build interactive ML interfaces without frontend code.
- Pre-built components for quick testing of ML models.
- Generates public links for real-time model interaction.
- Supports ML frameworks like TensorFlow, PyTorch, Hugging Face, and more.



BUILDING BLOCKS

Gradio - build fast ML webapps.

Let's write a chat function that responds **Yes** or **No** randomly.

Here's our chat function:

```
import random

def random_response(message, history):
    return random.choice(["Yes", "No"])
```

Now, we can plug this into `gr.ChatInterface()` and call the `.launch()` method to create the web interface:

```
import gradio as gr

gr.ChatInterface(random_response).launch()
```

A screenshot of a web application titled "Chatbot". It features a large text area for chat history, three buttons labeled "Retry", "Undo", and "Clear", a text input field with the placeholder "Type a message...", and an orange "Submit" button. At the bottom, it shows the text "gradio/chatinterface_random_response built with Gradio." and "Hosted on Spaces".

Chatbot

Retry Undo Clear

Type a message...

Submit

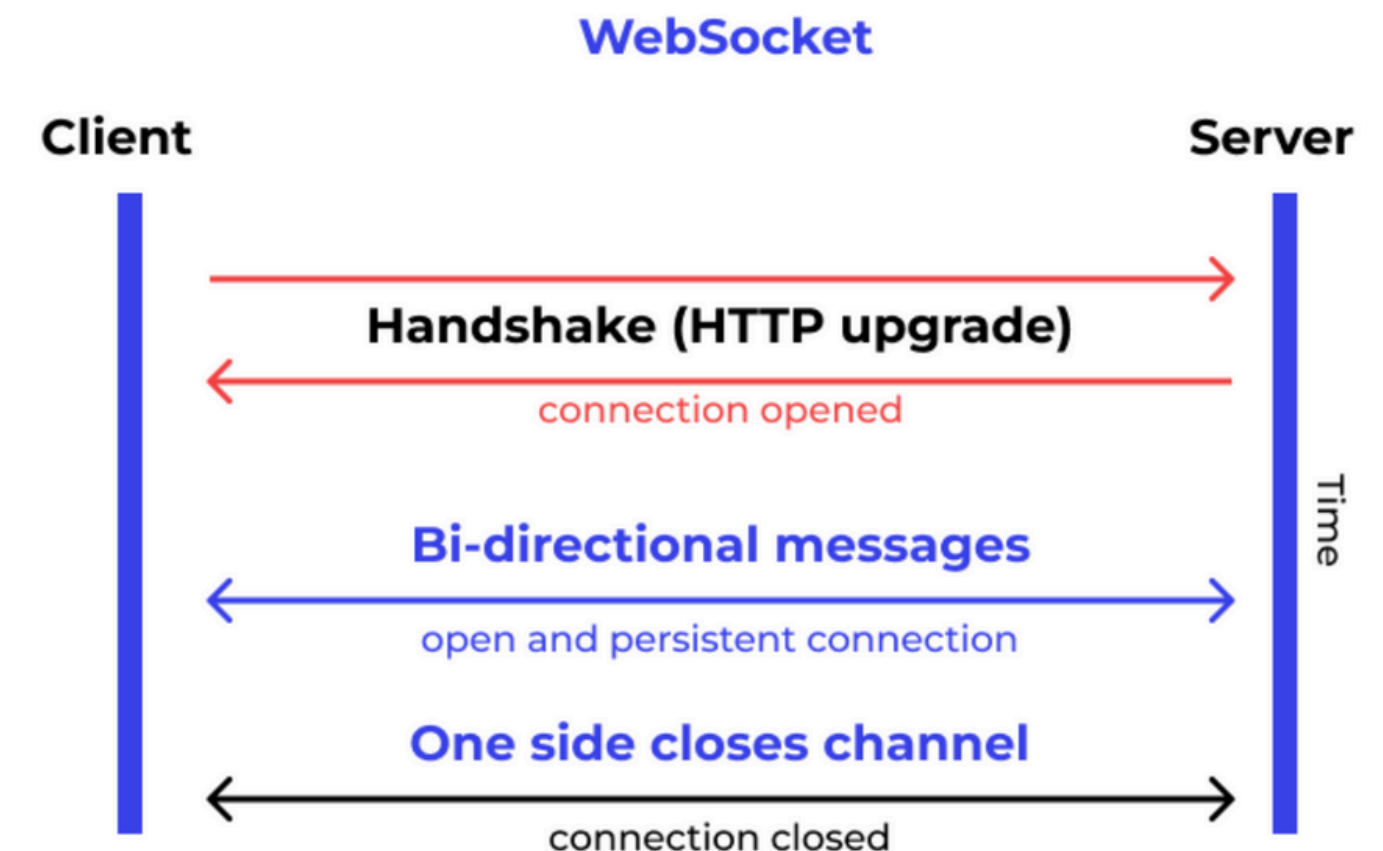
gradio/chatinterface_random_response built with Gradio. Hosted on Spaces

BUILDING BLOCKS

FastAPI Backend

- Asynchronous web framework optimized for building fast APIs.
- Simple syntax, leveraging Python type hints for automatic validation.
- Generates OpenAPI and Swagger documentation automatically.
- Supports async programming, WebSockets, and background tasks.

-> **we will build WebSocket API**

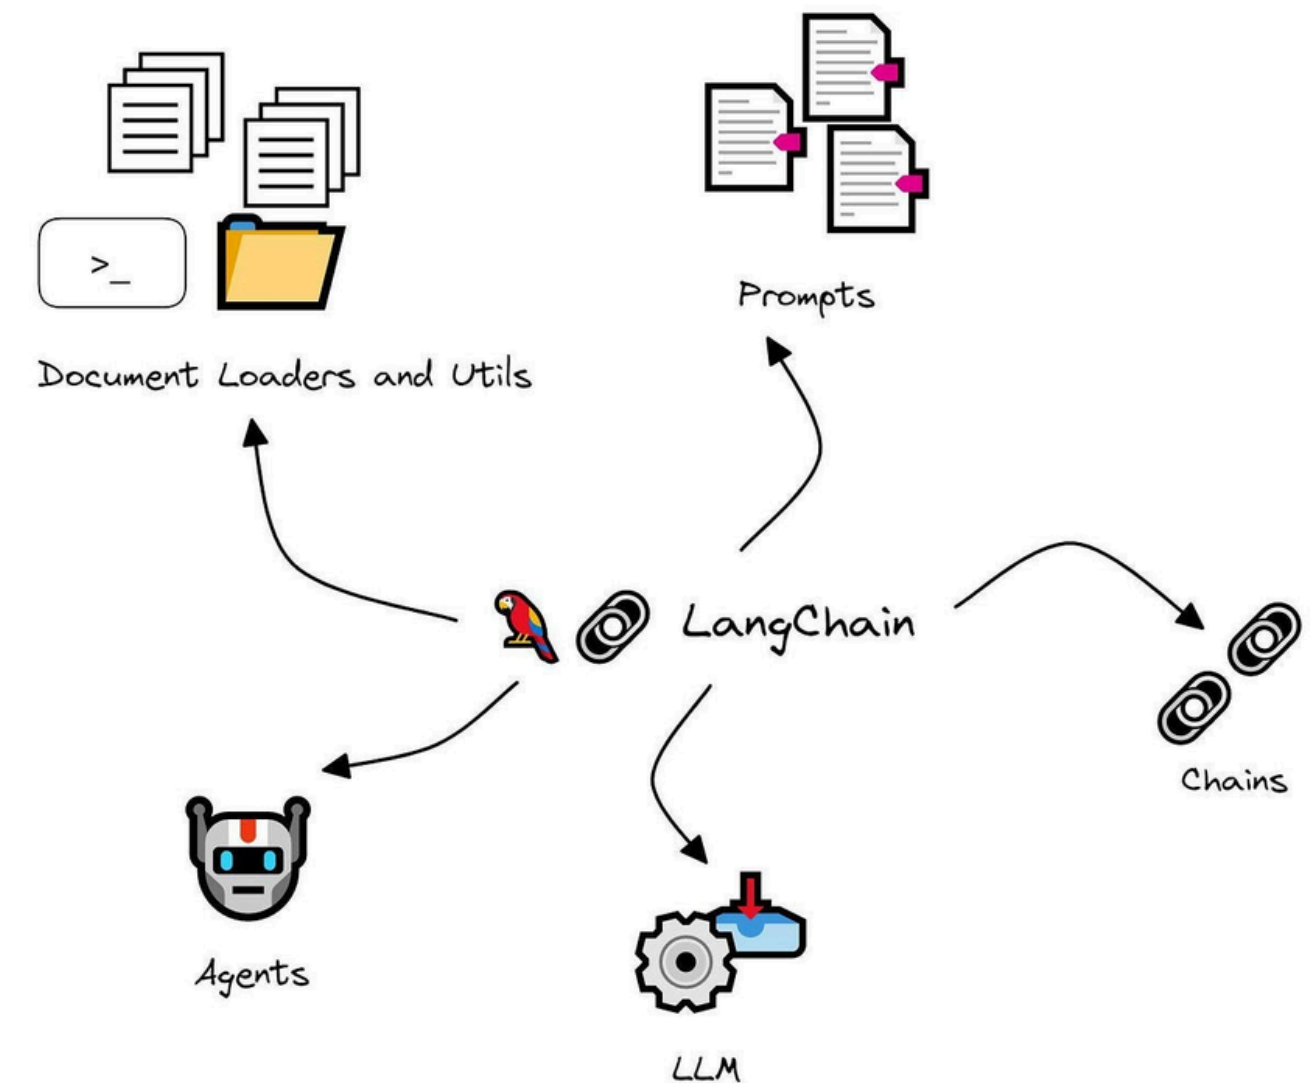


[HTTPS://WWW.WALLARM.COM/WHAT/A-SIMPLE-EXPLANATION-OF-WHAT-A-WEBSOCKET-IS](https://www.wallarm.com/what/a-simple-explanation-of-what-a-websocket-is)

BUILDING BLOCKS

RAG Chatbot with LangChain

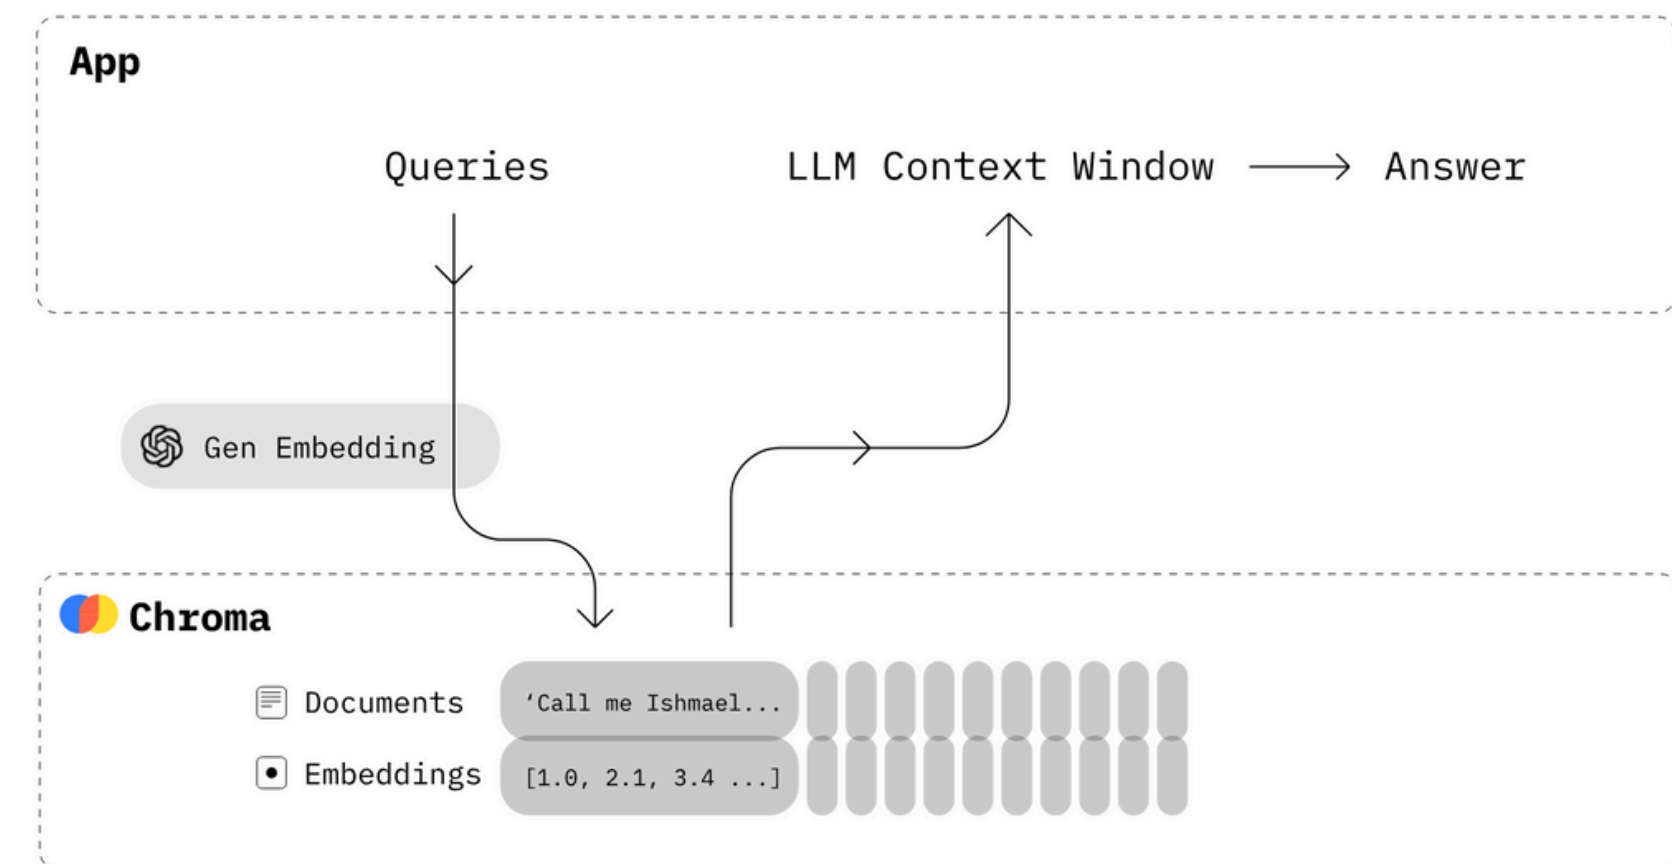
- Build LLM based apps
- Build complex pipelines by linking LLMs and tools.
- Supports APIs, databases, and custom logic for flexible workflows.
- Enables context persistence across multiple interactions.



BUILDING BLOCKS

Chroma as Vector Database

- Specialized for storing and querying high-dimensional embeddings.
- Designed to handle large-scale data efficiently.
- Works with popular ML frameworks like LangChain.
- Enables fast similarity searches for embeddings-based applications.




BUILDING BLOCKS

Ollama as LLM Runtime

- Run large language models on local machines efficiently.
- Keeps data local, ensuring better control over sensitive information.
- Designed for high-speed inference with minimal resource usage.
- Simple setup for running and experimenting with LLMs on your device.



```
nvidia@jao-60:/$ jetson-containers run $(autotag ollama) ollama run mistral
Namespace(packages=['ollama'], prefer=['local', 'registry', 'build'], disable=[''], user='dustynv',
output='/tmp/autotag', quiet=False, verbose=False)
-- L4T_VERSION=36.2.0 JETPACK_VERSION=6.0 CUDA_VERSION=12.2
-- Finding compatible container image for ['ollama']
cu122/ollama:r36.2.0
+ docker run --runtime nvidia -it --rm --network host --volume /tmp/argus_socket:/tmp/argus_socket
--volume /etc/encntune.conf:/etc/encntune.conf --volume /etc/nv_tegra_release:/etc/nv_tegra_release
--volume /tmp/nv_jetson_model:/tmp/nv_jetson_model --volume /var/run/dbus:/var/run/dbus --volume
/var/run/avahi-daemon/socket:/var/run/avahi-daemon/socket --volume /var/run/docker.sock:/var/run/docker.sock
--volume /mnt/NVME/jetson-containers/dev/data:/data --device /dev/snd --device /dev/vbus/usb --device /dev/video0
--device /dev/video1 cu122/ollama:r36.2.0 ollama run mistral
pulling manifest : ■
```



IT'S YOUR TURN